

Himanshu Jain

408-390-0164 | himanshu.jain@sjsu.edu | [LinkedIn](#) | [GitHub](#) | San Jose, CA

EDUCATION

San Jose State University

Master of Science in Software Engineering

San Jose, CA

Aug. 2025 - May 2027

National Institute of Technology, Warangal

Bachelor of Technology in Computer Science and Engineering

Warangal, India

Aug. 2019 - May 2023

EXPERIENCE

Software Engineering Intern | *Python, FastAPI, SQLAlchemy, PostgreSQL, RabbitMQ* June 2026 - August 2026

MalyTech

- Built **four asynchronous Python adapters** and a **non-blocking onboarding orchestrator** that connected a commercial bank's core to messaging, payments, identity-verification, and AML providers across five services.
- Engineered an **idempotent transfer submission engine** using atomic PostgreSQL reservations and payload hashing, ensuring eight concurrent requests produced **exactly one downstream provider call**.
- Delivered a **fail-closed AML gateway** that returned typed errors on provider faults, unified REST/JSON and SOAP/XML integrations, and enforced **credential isolation, PII redaction, and XML-injection defenses**.

Software Developer | *Java, Spring Boot, Kafka, Elasticsearch, OpenShift* July 2023 - June 2025

Amdocs Ltd.

- Cut idle Kafka topic sprawl by 70% and reclaimed ~80% of empty DLQ partitions** by consolidating 862 per-consumer dead-letter topics across 11 deployments into one shared DLQ per subdomain, delivered through a feature-flagged, zero-downtime Helm rollout.
- Enforced **tenant data isolation** by co-building a Java row-level entitlement library that translated JWT claims into Couchbase N1QL and Elasticsearch query predicates, pre-filtering rows and failing closed with HTTP 403.
- Shortened issue-resolution time by ~60%** by replacing multi-service log correlation with a single filtered query, building an append-only Kafka and Elasticsearch audit trail that indexed actor-attributed lifecycle events.
- Reduced production data footprint by ~40% and query latency by ~25%** by implementing policy-driven retention across three microservices, combining REST lifecycle APIs with Kafka archival under archive-before-delete ordering.
- Lowered manual end-to-end test effort by ~50% and cross-service defects by ~40%** by automating Postman and Newman suites with JSON-schema validation in Jenkins to catch contract drift at build time.
- Raised unit-test coverage to 90%** across four microservices with JUnit 5, Mockito, and WireMock, resolving ~500 SonarQube findings behind enforced build-time quality gates.

PROJECTS

ResumeFinetune - Agentic Job Intelligence Platform | *Python, Agentic AI, LLM Orchestration* Jul. 2026 - Present

- Engineered a deterministic Python job-intelligence pipeline with **ATS-specific resolvers** and typed eligibility gates, discovering 3,212 listings and resolving **2,019 job descriptions** to surface 164 qualifying roles from a seven-day discovery window.
- Cut mean LLM scoring movement from 0.67 to 0.20** using three-pass self-consistency and median aggregation, then enforced schema and row-coverage validation before **transactional score import**.
- Orchestrated a resumable, evidence-grounded tailoring workflow that constrained generated claims to **approved profile evidence** and routed drafts through deterministic validation, a separate LLM review, and human approval.
- Isolated LLMs behind a **least-authority, tool-free text boundary** with no repository or database access, while deterministic Python owned validation, trace capture, and **every filesystem and database write**.
- Enforced **fail-closed JD provenance** by binding content hashes, source and ATS metadata, and database verification or user attestation into each run, blocking unverified aggregator summaries and manifest drift.

Fake Review Detection on Yelp | *Python, PySpark, PyTorch, Transformers* Feb. 2026 - Apr. 2026

- Processed 608K Yelp reviews through a **PySpark ETL**, deriving 260K reviewer and 5K seller behavioral profiles into partitioned Delta Lake tables validated by automated quality checks.
- Fine-tuned a DeBERTa transformer on 423K labeled reviews using reviewer-disjoint data splits to prevent leakage, achieving **0.93 ROC-AUC and 0.84 macro-F1** on held-out test data.

PeerChat - Peer Discovery | *Python, WebSockets, Cryptography* Feb. 2026 - May 2026

- Designed the **event-driven membership layer** for a peer-to-peer chat system, building an append-only event log with deterministic snapshot projection and state transition guards to maintain consistent local peer rosters.
- Implemented the peer-to-peer network layer with **gossip dissemination** over an LRU-backed dedup cache, SWIM-inspired two-phase failure detection, and trust-on-first-use public-key exchange across a shared WebSocket transport.

TECHNICAL SKILLS

AI and Machine Learning: Agentic AI, LLM Orchestration, LLM Evaluation, Prompt Engineering, Structured Outputs, Human-in-the-Loop Systems

Languages: Java, Python, C++, SQL

Backend and Data: Spring Boot, FastAPI, PostgreSQL, Apache Kafka, RabbitMQ, Elasticsearch, PySpark, PyTorch

Developer Tools: Jenkins, Docker, Git/GitHub, Postman, OpenShift, SonarQube