

Himanshu Jain

408-390-0164 | himanshu.jain@sjsu.edu | [Linkedin](#) | [GitHub](#) | San Jose, CA

EDUCATION

San Jose State University

Master of Science in Software Engineering

San Jose, CA

Aug. 2025 - May 2027

National Institute of Technology, Warangal

Bachelor of Technology in Computer Science and Engineering

Warangal, India

Aug. 2019 - May 2023

EXPERIENCE

Software Engineering Intern | Python, FastAPI, SQLAlchemy, PostgreSQL, RabbitMQ

June 2026 - August 2026

MalyTech

- Built the **anti-corruption layer** of four adapter services between a commercial bank's core and providers for telecom messaging, real-time transfers, identity verification, and AML screening, then the **Core onboarding service** orchestrating three of those four through one non-blocking API, five asynchronous Python microservices in all.
- Engineered an **idempotent transfer submission engine** using atomic PostgreSQL reservations and payload hashing, ensuring eight concurrent requests produced **exactly one downstream provider call**.
- Designed a **fail-closed AML sanctions gateway** that returns a typed error instead of a false clear on any backend fault, serving both REST/JSON and legacy SOAP/XML through shared orchestration and hardened with **credential isolation**, **PII redaction**, and **XML-injection defenses**.

Software Developer | Java, Spring Boot, Kafka, Elasticsearch, OpenShift

July 2023 - June 2025

Amdocs Ltd.

- Cut idle Kafka topic sprawl by 70% and reclaimed ~80% of empty DLQ partitions** by consolidating 862 per-consumer dead-letter topics across 11 deployments into one shared DLQ per subdomain, delivered through a feature-flagged, zero-downtime Helm rollout.
- Enforced **tenant data isolation** by co-building a Java row-level entitlement library that translated JWT claims into Couchbase N1QL and Elasticsearch query predicates, pre-filtering rows and failing closed with HTTP 403.
- Shortened issue-resolution time by ~60%** by replacing multi-service log correlation with a single filtered query, building an append-only Kafka and Elasticsearch audit trail that indexed actor-attributed lifecycle events.
- Reduced production data footprint by ~40% and query latency by ~25%** by implementing policy-driven retention across three microservices, combining REST lifecycle APIs with Kafka archival under archive-before-delete ordering.
- Lowered manual end-to-end test effort by ~50% and cross-service defects by ~40%** by automating Postman and Newman suites with JSON-schema validation in Jenkins to catch contract drift at build time.

PROJECTS

ResumeFinetune - Agentic Job Intelligence Platform | Python, Agentic AI, LLM Orchestration

Jul. 2026 - Present

- Discovered 3,212 postings and resolved **2,019 job descriptions** in a 1-hour-43-minute run through a deterministic Python pipeline, surfacing 164 qualifying early-career roles from a seven-day discovery window.
- Cut mean LLM scoring drift from 0.67 to 0.20** by running three-pass self-consistency scoring with median aggregation, then validating output schemas before importing results into the production database.
- Engineered a resumable agentic workflow that turns a job description into an **evidence-grounded resume draft** through typed stage contracts, deterministic claim validation, an independent LLM review pass, and a human approval gate.

Fake Review Detection on Yelp | Python, PySpark, PyTorch, Transformers

Feb. 2026 - Apr. 2026

- Processed 608K Yelp reviews through a **PySpark ETL**, deriving 260K reviewer and 5K seller behavioral profiles into partitioned Delta Lake tables validated by automated quality checks.
- Fine-tuned a DeBERTa transformer on 423K labeled reviews using reviewer-disjoint data splits to prevent leakage, achieving **0.93 ROC-AUC and 0.84 macro-F1** on held-out test data.

Campus Marketplace | Java 21, Spring Boot, PostgreSQL

Sep. 2025 - Dec. 2025

- Built the backend for a campus marketplace as primary developer on a three-person team, delivering Java 21 and Spring Boot services over **Flyway-managed PostgreSQL** migrations.

TECHNICAL SKILLS

AI and Machine Learning: Agentic AI, LLM Orchestration, LLM Evaluation, Prompt Engineering, Structured Outputs, Human-in-the-Loop Systems

Languages: Java, Python, C++, SQL

Backend and Data: Spring Boot, FastAPI, PostgreSQL, Apache Kafka, RabbitMQ, Elasticsearch, PySpark, PyTorch

Developer Tools: Jenkins, Docker, Git/GitHub, Postman, OpenShift, SonarQube