

Himanshu Jain

408-390-0164 | himanshu.jain@sjsu.edu | [Linkedin](#) | [GitHub](#) | San Jose, CA

EDUCATION

San Jose State University

Master of Science in Software Engineering

San Jose, CA

Aug. 2025 - May 2027

National Institute of Technology, Warangal

Bachelor of Technology in Computer Science and Engineering

Warangal, India

Aug. 2019 - May 2023

EXPERIENCE

Software Engineering Intern | Python, FastAPI, SQLAlchemy, PostgreSQL, RabbitMQ

June 2026 - August 2026

MalyTech

- Built the **anti-corruption layer** of four adapter services between a commercial bank's core and providers for telecom messaging, real-time transfers, identity verification, and AML screening, then the **Core onboarding service** orchestrating three of those four through one non-blocking API, five asynchronous Python microservices in all.
- Engineered an **idempotent transfer submission engine** using atomic PostgreSQL reservations and payload hashing, ensuring eight concurrent requests produced **exactly one downstream provider call**.
- Designed a **fail-closed AML sanctions gateway** that errors on backend faults or invalid responses instead of returning a false clear, serving REST/JSON and legacy SOAP/XML from shared orchestration; hardened all four services with **credential isolation, recursive PII redaction, and XML-injection resistance**.

Software Developer | Java, Spring Boot, Kafka, Elasticsearch, OpenShift

July 2023 - June 2025

Amdocs Ltd.

- Cut idle Kafka topic sprawl by 70% and reclaimed ~80% of empty DLQ partitions** by consolidating 862 per-consumer dead-letter topics across 11 deployments into one shared DLQ per subdomain, delivered behind a feature-flagged **Helm binding function**.
- Cut issue-resolution time by ~60%** by replacing multi-service log correlation with a single filtered query, building an **append-only Kafka and Elasticsearch audit trail** that indexed actor-attributed lifecycle events with correlation IDs.
- Reduced production data footprint by ~40% and query latency by ~25%** by implementing policy-driven retention across three microservices, combining REST lifecycle APIs and **Kafka archival with archive-before-delete ordering**.
- Reduced manual end-to-end effort by ~50% and cross-service defects by ~40%** by automating Postman/Newman suites with **JSON-schema validation in Jenkins** to catch contract drift at build time.
- Enforced **tenant data isolation** by co-building a Java row-level entitlement library that translated JWT claims into **Couchbase N1QL and Elasticsearch query predicates**, pre-filtering rows and failing closed with HTTP 403.
- Built and maintained the **Java and Spring Boot microservices** behind the Order Management domain across Catalog, Shopping Cart, Proposal, and Agreement, automating builds, tests, and deployments to OpenShift with a **Jenkins CI/CD pipeline**.
- Raised unit-test coverage to 90%** across four microservices using **JUnit 5, Mockito, and WireMock**, resolving ~500 SonarQube findings behind enforced deployment quality gates.

PROJECTS

ResumeFinetune - Job Search and Resume Tailoring Platform | Python, SQLite, LLM APIs, pytest, LaTeXJul. 2026 - Present

- Surfaced **164 qualifying US early-career roles** from a seven-day discovery window by building a deterministic Python pipeline that discovered 3,212 listings, inserted 3,176 records, and resolved **2,019 current and backlog job descriptions** in 1 hour 43 minutes.
- Implemented a **resumable agentic AI workflow** for resume tailoring across requirement extraction, evidence selection, bounded alignment, deterministic lint, independent LLM evaluation, **one-page LaTeX rendering**, and human review.

PeerChat - Peer Discovery | Python, WebSockets, cryptography, pytest

Feb. 2026 - May 2026

- Designed the event-driven membership layer for a **peer-to-peer chat system**, building an **append-only event log with deterministic snapshot projection** and state transition guards to maintain consistent local peer rosters.
- Resolved a cross-machine discovery failure by **consolidating peer discovery onto a shared WebSocket connection**, eliminating a redundant TCP/crypto stack and leaving **one port and one retry queue**.

Campus Marketplace | Java 21, Spring Boot, PostgreSQL, AWS S3

Sep. 2025 - Dec. 2025

- Built the backend for a campus peer-to-peer marketplace as primary developer on a three-person team, delivering a domain-organized **Java 21 and Spring Boot** service over **Flyway-controlled PostgreSQL** schema migrations.

TECHNICAL SKILLS

AI and Machine Learning: LLM Integration, Prompt Engineering, AI Agent Design, Structured Output Parsing | **Languages:** Java, Python, C++, SQL | **Frameworks:** Spring Boot, Spring MVC (REST), FastAPI, JUnit, Mockito | **Libraries:** Apache Kafka, RabbitMQ, Elasticsearch, PyTorch, scikit-learn, XGBoost, PySpark | **APIs and Standards:** REST, JSON, OpenAPI/Swagger, OAuth2/JWT, HTTP | **Developer Tools:** Jenkins, Maven, Docker, Git/GitHub, Bitbucket, Postman, OpenShift (OCP), SonarQube, IntelliJ | **Databases:** PostgreSQL, Couchbase (SDK, N1QL), SQL (RDBMS)